

- Heap:**
- implemented as a complete binary tree in an array
 - for node at index i
 - * left child is at $2i$
 - * right child is at $2i + 1$
 - * parent is at $\lfloor \frac{i}{2} \rfloor$ (floor)
 - insert: put at the last node and up-heap if needed. $O(\log(n))$
 - remove min: replace the root node with the last node, and down-heap if needed. $O(\log(n))$
 - up-heap: compare a node with it's parent. swap if needed. Repeat until you no longer need to swap, or at the root node
 - * down-heap: same as up-heap, except going down. choice of left or right child is arbitrary.
 - merge two heaps (of equal height) with an extra element e : just connect e as the new root node, then down-heap if needed
 - bottom-up heap construction:
 - * treat array as complete binary tree
 - * down-heap starting at second-to-bottom row (bottom row is leaves, therefore satisfies heap property)
 - * continue down-heap until you get up to the root node
 - * not every node can be down-heaped the full h , because they started at different points
 - this is why it's $O(n)$

- Priority Queue:**
- Comparison:

implementation	insert	removeMin
unsorted list	$O(1)$	$O(n)$
sorted list	$O(n)$	$O(1)$
heap	$O(\log(n))$	$O(\log(n))$
 - priority queue sort: make the input into a priority queue (in-place) then remove each from the queue back into the list
 - * complexity is $O(n * \text{insert} + n * \text{removeMin})$
 - PQ sorts: selection, insertion, heap
 - * selection: unordered list
 - * insertion: ordered list
 - * heap: heap (duh)

- Set:**
- ordered
 - union: $A \cup B$: all elements in either A or B
 - intersection: $A \cap B$: all elements in both A and B
 - subtraction (relative complement): $A \setminus B$: all elements in A and not in B
 - can be easily implemented using a binary tree

- Quick Select:**
- algorithm to find the k th smallest element
 - algorithm that's the same as quicksort, except it only recurses onto the section containing the element we're looking for
 - recurrence relation: $T(n) = O(n) + T(\frac{n}{2})$
 - * $O(n)$ for the partitioning step
 - * this solves to $O(2n)$ which is $O(n)$

- Sorting:**
- slow sorts:**
- selection, insertion, bubble
 - all $O(n^2)$ average
 - insertion and bubble are $O(n)$ best-case
 - selection is $O(n^2)$ always

- Heap Sort:**
- $O(n \log(n))$ all cases
 - in-place
 - can use the bottom-up heap building to be faster, but the removal stage still limits it to $O(n \log(n))$
 - non-recursive

- Merge Sort:**
- $T(n) = 2T(\frac{n}{2}) + O(n)$
 - * splitting is $O(1)$, merging is $O(n)$
 - always $O(n \log(n))$
 - not in-place

- recursive
- Quick Sort:**
- $T(n) = 2T(\frac{n}{2}) + O(n)$
 - recursive
 - pivot is randomly selected
 - $O(n \log(n))$ average and best, $O(n^2)$ worst
 - * worst-case: pick worst element every time. This makes it $T(n) = T(1) + T(n + 1) + O(n)$.
 - * Unlikely for random numbers

- Graphs:**
- path: sequence alternating vertexes and edges. must begin and end on a vertex.
 - * simple path has all vertexes and edges unique (does not cross itself)
 - cycle: a loop of nodes
 - * simple if it doesn't intersect itself (except beginning/end)
 - * non-simple otherwise
 - total in-degree of a graph is equal to it's total out degree
 - total degree of graph is double the number of edges
 - connected graph: there is a path between every pair of vertexes
 - tree: connected graph with no cycles
 - spanning tree: tree that includes all vertexes in graph
 - TODO expound this

	Edge List	Adjacency List	Adjacency Matrix
space	$O(n + m)$	$O(n + m)$	$O(n^2)$
endVertexes()	$O(1)$	$O(1)$	$O(1)$
opposite()			
incidentOn(v)			
v.incidentEdges()	$O(m)$	$O(deg(v))$	$O(n)$
v.adjacentTo()	$O(m)$	$O(\min(deg(v), deg(w)))$	$O(n)$
insertEdge(u,v,w)	$O(1)$	$O(1)$	$O(1)$
eraseEdge(e)			
insertVertex(x)	$O(1)$	$O(1)$	$O(n^2)$
eraseVertex(v)	$O(m)$	$O(deg(v))$	$O(n^2)$

- adjacency matrix is usually a bad choice
- DFS: Depth First Search:**
- visits each child node before visiting adjacent nodes
 - can be used for maze traversal
 - * each position is a vertex, edges are places you can get to
 - similar to preorder tree traversal
 - labels vertexes visited or not, labels edges as discovery or back edges
 - * no cross edges because those would have been discovery edges
 - $O((m + n))$ for adjacency list?
 - DFS and BFS are both good for finding connected components

```
DFS(G: (V,E), s: starting v):
    v.label = VISITED
    for e in v.incidentEdges():
        if e.label == UNEXPLORED:
            w = e.opposite(v)
            if w.label = UNEXPLORED:
                e.label = DISCOVERY
                DFS(G, w)
            else:
                e.label = BACK
// then repeat for other connected components
```

- BFS: Breadth First Search:**
- visits all nodes in order of their distance from the starting node. (distance in hops, not counting edges)
 - labels vertexes to keep track of whether visited
 - labels edges as discovery or cross edges
 - * no back edges because those would have been discovery edges (assuming undirected)
 - uses a (non-priority) queue to keep track of the edges to check next
 - $O((m + n))$ for adjacency list

BFS($G: (V,E)$, s : starting v):

```
Q.enqueue(s)
while !Q.empty():
    v = Q.dequeue()
    v.label = VISITED
    for e in v.incidentEdges():
        if e.label == UNEXPLORED:
            w = e.opposite(v)
            if w.label == UNEXPLORED:
                e.label = DISCOVERY
                w.label = VISITED
                Q.enqueue(w)
            else:
                e.label = CROSS
// then repeat for other connected components
```

MST: Minimum Spanning Tree:

- minimum total weight spanning tree
- partition property: if you partition the MST into two subsets, there must be exactly one edge connecting the two, and it must be the minimum possible edge connecting the two subsets

Prim-Jarnik's Algorithm:

- $O((m+n)\log(n))$ for adjacency list
- start at a given (or arbitrary) node as our MST
- put all the vertexes in a PQ keyed with their shortest edge connecting them to the MST
- add the closest vertex v to the MST
- update the vertexes adjacent to v with their new distance (use $PQ.replaceKey()$)
- repeat until the PQ is empty

```
Prim_Jarnik(G: (V,E), s):
    for each v in V:
        D[v] = inf, P[v] = NULL
        PQ.add(v, key=D[v])
    while (!PQ.empty()): // O(n)
        u = PQ.removeMin()
        for e in u.edges(): // O(m)
            z = e.opposite(u)
            if e.weight() < D[z]:
                D[z] = e.weight(); P[z] = u
            PQ.replaceKey(z, D[z]) // O(log(n))
```

Kruskal's Algorithm:

- $O((m+n)\log(n))$ for adjacency list
- initialize a PQ with all edges keyed by weight
- make clouds of mini MST's by adding each minimum edges
 - * adding edge joining non-cloud vertex to cloud is easy
 - * if an edge connects two vertexes in the same cloud, ignore it
 - * if an edge connects two vertexes in different clouds, add it and merge the clouds (merging these clouds is complex depending on the implementation)
- does not start at any particular node
- cluster merging: $merge(u,v)$ is $O(\min(|C_u|, |C_v|))$. We assume that merging doubles the size of the sets, therefore we preform $\max \log(n)$ merges
- complexity:
 - * PQ: m removals: $O(m\log(n))$; cluster merges: $O(n\log(n))$
 - * total: $O((m+n)\log(n))$

```
kurskals(G: (V,E)):
    T = (V, NULL) // all vertexes, no edges
    for v in V: { define cluster C(v) = {v} }
    for e in E: { PQ.add(e, key=e.weight()) }
    while ( P.size() < V.size() - 1 ): // O(m)
        (u,v) = PQ.removeMin() // O(log(n))
        if C(u) != C(v): // nodes from different clusters
            T.insertEdge(u,v)
            MergeClusters(C(u), C(v) )
            // ~- O(min(|C(u)|, |C(v)|))
```

SSSP: Single Source Shortest Path:

- SSSP is a spanning tree where the path from every node to the root is the shortest possible path.
 - * not necessarily the same thing as the MST
- a subpath of a shortest path is itself a shortest path
- if there is no path between two vertexes, we generally represent the path length as ∞
- **Dijkstra's Algorithm**
 - * $O((m+n)\log(n))$ for adjacency list
 - $(m+n)$ because it must look at every node
 - $\log(n)$ for the PQ operations

- * assumes: graph is connected, all edges are undirected, all weights are ≥ 0
- * is a greedy algorithm
- * very similar to Prim-Jarnik's Algorithm, main difference is that we care about the distance to the root node, not the distance to the cloud
- * store all distances in a map keyed with the vertex; $\text{map}<\text{dist}, \text{vertex}> D[]$
- * also store vertexes in a PQ, keyed with their total distance from the root (also stored in $D[]$)
- * edge relaxation:

- for vertex v_0 not yet in the cloud: check if edge (v_b, v_0) provides a shorter path than the current edge (v_a, v_0)
- if $D[v_b] + (v_b, v_0).weight < D[v_a] + (v_a, v_0).weight$: use the new edge (v_b, v_0) to connect v_0 to the cloud, and update $D[v_0]$ accordingly
- * does not work for negative edge weights because it is greedy, doesn't go back and check for the ways negative weighs could change things

```
dijkstras(G: (V,E) P: ParentMap, s: start vertex):
    D[v] = infinity for each v in V
    D[s] = 0
    P[s] = NULL
    Q = PQ of all v in V keyed with D[v]
    while !Q.empty(): // O(n)
        u = Q.removeMin() // O(log(n))
        for e in u.edges(): // relax each adj. edge
            z = e.opposite(u)
            if D[u] + e.weight() < D[z]:
                D[z] = D[u] + e.weight()
                Q.updateKey(z) // O(log(n))
                P[z] = u // update the parent of this node
```

• Bellman-Ford Algorithm

- * $O(nm)$
- * works for negative edge-weights (therefore must assume directed graph, otherwise there are negative weight cycles)
- * doesn't work if there are negative weight cycles, but can be extended to detect them
- * iteration i finds all shortest paths of length i , therefore the last iteration finds the maximum length shortest-path, of length $|V| - 1$
 - detect negative weight cycles: relax again at the end. If any edge can be relaxed, there is a shortest path with length $|V|$, and therefore there must be a negative weight cycle

```
bellman_ford(G: (V,E), s: starting vertex, P):
    D[v] = infinity for each v in V
    D[s] = 0; P[s] = NULL
    for i = 1:(V.size() - 1): // O(n)
        for each e in E: // O(m)
            // relax edge
            u = e.source(); z = e.target()
            if D[u] + e.weight() < D[z]:
                D[z] = D[u] + e.weight()
                P[z] = u
```

Directed Graphs:

- also just called digraph
- graph is **strongly connected** if every vertex can be reached from every other vertex
- strongly connected components: subsets of a graph which are themselves strongly connected.
- determine if graph G is strongly connected:
 - * do DFS on G . if there are any nodes not visited, graph is not strongly connected.
 - * $G' = G$ with all directed edges reversed
 - * do DFS on G' . if there are any nodes not visited, graph is not strongly connected.
 - * otherwise, graph is strongly connected
- Directed Acyclic Graph: directed graph with no directed cycles
- DFS and BFS make sense on a digraph. MST doesn't really make as much sense on a digraph