

Batch OS

- jobs are submitted in batches
- just used DMA for IO
- basically FIFO, but with using DMA for loading P2 while P1 is still running

Time Sharing OS

- CPU cycles through jobs in defined order
- basically round robin

Exception Control Flow

- exception: transfer of control from the user program to the OS in response to an event
- event: interrupt or system call or something
- CPU checks for interrupts before every instruction
- **Asynchronous Exception:** caused by something external to the processor. (e.g. not a system call)
 - * e.g. IO interrupts, hard/soft reset
- **Interrupt Vector Table:** stored in CPU, has function pointers at indexes identified by interrupt codes
 - * one slot for every possible type of interrupt
 - * first thing handler does is disable interrupts, then re-enable when done
- **Synchronous exceptions:** caused by executing an instruction (internal to CPU)
 - * Trap: intentional; system calls, breakpoints
 - returns code to next userland instruction
 - * Faults: unintentional. maybe recoverable; generally retry or abort. e.g. page fault, floating point exception
 - * Abort: unintentional; unrecoverable
- RTU: Return To User

Dual Mode

- CPU supports two modes: kernel mode and user mode
 - * keeps a flag register depending on current mode
- user mode is disallowed some dangerous instructions and only allowed it's mapped memory regions
- kernel mode gets full control over everything
- must switch between the two, which has overhead
- Memory Protection: uses MMU hardware
 - * kernel has no memory protection
- hardware timer: fires interrupts at predefined intervals
 - * kernel will set timed interrupts to switch control back to the kernel from the user program periodically
 - * setting/clearing timers is a privileged instruction
- user→kernel: on interrupt, syscall
- kernel→user: end of interrupt handler, or at proc start

Unix Processes

- process: instance of running program
- has it's own private address space
- **Zombie Process:** when a child process exits, the kernel must keep it's memory around just in case the parent needs to know if it exited cleanly
 - * this child process is a zombie process.
 - * call `wait()` to avoid this
 - * reaping is the process of killing zombies

Process Lifetime/States

- all this stuff looks like continuous running to the process
- **Ready:** scheduled to be run
- **Running:** currently running
- **Blocked:** waiting on something. (IO or other wait)
- transitions:
 - * newly created processes go to ready; after exit, leaves running position
 - * ready→running: dispatch
 - * running→ready: preempt
 - also happens when scheduling timer expires
 - * running→blocked: triggers IO event or wait
 - * blocked→ready: IO event or wait finishes
- only one process is in running state per CPU core
- Additional states: (managed by memory scheduler)
 - * **Suspended Ready:** ready to run, but swapped out of main memory due to memory constraints. ; get here

from start or ready ; leave to ready when ready

- * **Suspended Blocked:** swapped out of main memory, and also waiting on some external event ; get here from blocked ; leave to suspended ready or blocked
- **Process Control Block (PCB)**
 - * contains process id, state, saved registers, memory maps, child processes, owned resources, file descriptors
 - * whenever a context switch happens, the previous process state is stored in and restored from the PCB
- **Job Queue:** set of all processes in a system
- **Ready Queue:** in main memory, waiting to be run
- **Device Queues:** waiting for IO by device

Scheduling

- **Latency:** how long it takes for the job to complete
- **Throughput:** # of jobs can be completed per unit time
- **Overhead:** how much work the scheduler must do
- **Fairness:** do different users/process types/other factors influence scheduling priority?
- **Predictability:** also consistency
- **Preemptive Scheduler:** one that takes control away from a process (usually through timed interrupts)
- **Work Conserving:** never leave a CPU idle
- **Time Quantum:** length of scheduler interrupt timer
- **Waiting Time:** time spend in not running
- **Service (Execution) Time:** how long the job is running
- **Response (Completion) Time:** wall clock time process takes ; response = waiting + ready
- **Long-Term Scheduler:** (Job Scheduler)
 - * selects which job to brought into the ready queue
 - * should select a good mix of CPU and IO bound processes
- **Short-Term Scheduler:** (CPU Scheduler)
 - * selects what to run from the ready queue
 - * runs every time quantum, so it shall be fast
- **Medium-Time Scheduler:** swaps out programs from main memory (suspends them)
- **First In First Out (FIFO), or FCFS:**
 - * works well on one really long task with many small tasks, if the long task comes first (e.g. servers)
- **Shortest Remaining Time First: (SRTF)** (or Shortest Job First (SJF))
 - * the ideal scheduler in most cases
 - advantage best for average response time
 - disadvantage: longer tasks can get starved
 - disadvantage: high variance on average response time
 - * not really possible because we can't know how long a job will take ahead of time
 - can be approximated by guessing the next CPU burst
 - * if all tasks are the same length, runs them all in order, no preemption
 - * shortest task starts and finishes first
- **Round Robin:**
 - * each task gets run for one time quantum
 - choice of time quantum is critical
 - * compromise between SJF and FIFO
 - * always fair
 - * time quantum too small: too much overhead
 - * time quantum approaches ∞ : equivalent to FIFO
- **Multi-Level Feedback Scheduling:**
 - * have multiple ready queues
 - sub-queue are prioritized
 - * processes start in highest priority queue, then are moved to lower queue after using CPU too much
 - * must schedule between queues:
 - fixed priority: highest priority, then lower (not fair)
 - time slicing: fixed percentage of time to each queue.
 - * approximates SRTF: CPU-bound suffers, lets IO bound stay near top
 - * example: putting in superfluous IO waits
- utilization approaches 100%, latency and throughput suffer