

CSCE 441 Computer Graphics

scan conversion of lines

- horizontal, vertical lines are easy
- for general lines, assume $0 < slope < 1$ (flat to diagonal)
 - you can transform any line to fit this
- naive algorithm would just use floating point and round off
 - floating point is sometimes slow (especially back when not every computer did it in hardware)
- slope from two points:

$$m = \frac{y_H - y_L}{x_H - x_L} a$$

- $s \frac{a}{b}$
- intercept from two points: $b = y_L - m * x_L$
- **Simple Algorithm**
 - start from (xL, yL) and draw to (xH, yH)
 - * where $xL < xH$
 - ```
def draw_line(xL, yL, xH, yH):
 x, y = (xL, yL)
 for i in range(0, xH - xL):
 draw_pixel(x, round(y))
 x = x + 1
 y = m * x + b # simplifies to
 y = y + m
```

- problem: uses floating point math
- problem: rounding

- **Midpoint Algorithm**
  - given a point, we just need to know whether we will move right or up and right on the next step (N or NE)
  - we can simplify this to whether the actual line travels above or below the point  $(x + 1, y + 1/2)$ 
    - \* so we derive formula from  $y = m * x + b$
  - formula:  $f(x, y) = c * x + d * y + e$ 
    - \*  $c = yL - yH$
    - \*  $d = xH - xL$
    - \*  $e = b * (xL - xH)$
    - \*  $f(x, y) = 0$ :  $(x, y)$  is on the line
    - \*  $f(x, y) < 0$ :  $(x, y)$  below line
    - \*  $f(x, y) > 0$ :  $(x, y)$  above line
  - don't want to recalculate formula at every step, so do it iteratively
    - \* that is, use  $f(x + 1, y + 1/2)$  to calculate  $f(x + 2, y + 1/2)$  or  $f(x + 2, y + 3/2)$  depending on right or up-right choice last time
  - went right last time, now calculate  $f(x + 2, y + 1/2)$

- \*  $f(x + 2, y + 1/2) = c + f(x + 1, y + 1/2)$
- went up-right last time, now calculate  $f(x + 2, y + 1/2)$ 
  - \*  $f(x + 2, y + 3/2) = c + d + f(x + 1, y + 1/2)$
- starting value:  $f(x + 1, y + 1/2) = f(xL, yL) + c + (1/2)d = c + (1/2)d$ 
  - \* we can eliminate  $f(xL, yL)$  because we know it is on the line
  - \* furthermore, we can use  $f(x + 1, y + 1/2) = 2 * c + d$  because multiplying by 2 does not change the sign of  $f$ . Also, this saves an expensive division
- full algorithm:
 

```
def midpoint_algorithm_line(xL, yL, xH, yH):
 x = xL
 y = yL
 d = xH - xL
 c = yL - yH
 sum = 2*c + d
 draw_pixel(x,y)
 while x < xH:
 if sum < 0:
 sum += 2*d
 y += 1
 x += 1
 sum += 2*c
 draw_pixel(x,y)
```
- pro:
  - \* only integer operations
  - \* extends to other kinds of shapes, just need formula to tell if inside/outside shape (called implicit formula)
- same as Bresenham's algorithm (more common algorithm)

## scan conversion of polygons

- to deal with overlap, we do not draw the top and right of a polygon
  - this means artifacts are possible. This doesn't really matter since pixels are very small
- rectangles (aligned with axes) are easy
- scan line: one row of pixels
- general polygons: basic idea (**scanline method**)
  - intersect scan lines with edges of polygon
  - this means you must keep track of which edges intersect with which scan lines
    - \* this is easy to do: just look at the y coordinate
  - consecutive scan lines will usually intersect with a similar set of edges
    - \* so we can use coherence to speed stuff up
  - we can throw out horizontal lines. They are implicitly represented by

- start and end, connecting to the other edges
- data structures
  - \* edge: maxY, currentX, xIncr (increment)
    - calculate these from the two points
    - xIncr is inverse of slope, but you can't calculate the slope and invert it, because divide by 0
    - maxY: y value of higher point
    - currentX: x value of lower point
  - \* active edge table
    - has entry for every scanline on the screen
    - initialize with edges by indexing by minY of edge
  - \* active edge list
    - stores edges that intersect with the current scan line being processed
    - edges must always be sorted by current x value
- at each step of the algorithm, you must update the active edge list
  - \* remove edges where maxY is less than or equal to the current scan line
    - less or equal because we don't draw the top and right of the polygon
  - \* add edges from the current scan line to the edge list
  - \* sort all edges by currentX
- then draw the scan line
  - \* take pairs of edges and fill in between their currentX values
    - do not include the right point (because we don't draw the top and right of the polygon)
  - \* if you ever have an odd number of edges in the active edge list, you made a mistake
- disadvantages
  - \* does not handle long, thin polygons well
  - \* incremental updates are not suitable for massively parallel GPUs
- boundary fill
  - draw the boundary of the polygon, then fill in interior
    - \* fill in interior wherever it is not the same color as you are drawing
  - need to be sure filling can't escape out from an edge or corner
  - need to be able to choose arbitrary interior point to start from
- flood fill
  - starting at point, recursively replace one color with another
  - paint bucket tool

## openGL data CPU to GPU

- openGL can accept data various ways, with different speed impacts
- speed depends on driver implementation

- GPUs only render triangles, and triangles usually share vertexes with other triangles, so saving lots of bandwidth is possible
- fastest is usually vertex buffer objects?
  - stores data directly on GPU?

## clipping lines

- it's not really possible to draw things that are outside of the viewing area
- clipping points is easy (when comparing to rectangular window)
- clipping lines:
  - if both end points are inside window, draw it
- window intersection method:
  - if either or both is outside, intersect line with each window border in sequence
  - $(x_1, y_1), (x_2, y_2)$  intersect with vertical edge at  $x_{right}$ :  
 $y_{intersect} = y_1 + m * (x_{right} - x_1)$ , where  $m = (y_2 - y_1) / (x_2 - x_1)$
  - $(x_1, y_1), (x_2, y_2)$  intersect with horizontal edge at  $y_{bottom}$ :  
 $x_{intersect} = x_1 + (y_{bottom} - y_1) / m$ , where  $m = (y_2 - y_1) / (x_2 - x_1)$
  - all these intersections are costly to compute
    - \* we would like to efficiently handle trivial accepts and trivial rejects
- cohen-sutherland algorithm
  - classify two points  $p_1, p_2$  using 4-bit codes **c0** and **c1**
  - if **c0** & **c1** != 0: trivial reject
    - \* bitwise AND
    - \* both points are outside one of the boundaries
  - if **c0** | **c1** == 0: trivial accept
    - \* bitwise OR
    - \* none of the coordinates of either point is outside any boundary  
=> line is entirely within window
  - otherwise split line until it is a trivial case
  - bits: | **top** | **bottom** | **right** | **left**
    - \* doesn't matter as long as you're consistent? TODO
    - \* you can determine each of these by just comparing one coordinate with the axes
    - \* thus the comparison is fast
  - disadvantages
    - \* repeated clipping is expensive
  - advantages
    - \* considers all possible trivial accept/reject
- laing-barsky algorithm
  - use parametric form of line for clipping
    - \* means that lines are oriented (have a direction)
  - need to classify lines as moving into or out of the window

- since lines are parametric, we will be finding the parameter value of the intersection
  - \* we can put that back into the formula to get the actual point
- parametric lines
  - \*  $x(t) = x_0 + (x_1 - x_0) * t$
  - \*  $y(t) = y_0 + (y_1 - y_0) * t$
  - \*  $0 \leq t \leq 1$
  - \* solve 2d matrix to intersect lines:
 
$$\begin{bmatrix} x_1 - x_0 & x_2 - x_3 \end{bmatrix} \begin{bmatrix} t \\ s \end{bmatrix} = \begin{bmatrix} x_2 - x_0 \\ y_2 - y_0 \end{bmatrix}$$
- algorithm:
  - \* start with  $t$  on range  $[0, 1]$ 
    - this is  $t_{min}, t_{max}$
  - \* iteratively intersect each line with each edge
    - find intersection at  $t$
    - if line is moving in to out:  $t_{max} = \min(t_{max}, t)$
    - else:  $t_{min} = \max(t_{min}, t)$
    - if  $t_{min} > t_{max}$ : reject line
- moving out vs moving in can be determined by looking at coordinates
  - \* different for each boundary
  - \* e.g. for right boundary,  $x_1 < x_2$  is moving in
  - \* does not depend on where either point is, or whether either point is inside/outside window boundary, just relative positions of the points
- disadvantages
  - \* does not consider trivial accept/reject
- advantages
  - \* computation of  $(x, y)$  is done only once at the end
  - \* computation of parametric intersections is fast (only one division)
- note: clipping line and then rounding to integer coordinates may not produce the correct result, due to round-off error
  - can account for this by calculating sum for use in midpoint algorithm

## clipping polygons

- clipping a polygon can change the number of sides it has
  - minimum number of sides is 3 (triangle)
  - maximum number of sides is  $2n + 1$ ? TODO
  - e.g. maximum number of sides of triangle after clipping is 7 sides
- when clipping convex polygons, you could end up with multiple polygons
- sutherland-hodgman clipping
  - clip polygon vs each edge of window individually
    - \* thus can handle non-rectangular window, as long as the window is convex

- is not guaranteed to handle convex polygons correctly
  - \* does not split into multiple polygons
  - \* but usually looks about right
- output is mixture of old/new vertexes
  - \* will be exactly old vertexes if polygon was entirely inside the window
  - \* will be only new vertexes if all vertexes were outside the window (but not necessarily all edges)
- process each side of the rectangular window separately
  - \* and also, process each edge in polygon iteratively
- 4 cases for an edge from S to E:
  - \* S and E both outside: no output
  - \* S and E both inside: output only E
  - \* S inside, E outside: compute intersection with border, and output that
  - \* S outside, E inside: output intersection with border, and output E
- output of one intersection is used as input for next intersection
  - \* you can kind of do these in parallel, with the partial output from the previous stage
    - pipeline
  - \* then you need a end-of-polygon marker, and you need to use that along with the first edge to make the last edge
- weiler-atherton algorithm
  - general intersection between any two kinds of polygons
  - handles non-convex polygons
    - \* thus can output more than one polygon for a single input polygon
  - not as efficient as sutherland-hodgman
    - \* all those intersections are expensive
    - \* difficult to parallelize
  - algorithm
    - \* start at point on polygon
    - \* follow polygon edges counterclockwise until an edge crosses out of the window
    - \* follow window edges from the intersection point until the polygon intersects again
    - \* now that part is a polygon. Go back to the first intersection point and follow the polygon until it re-enters the window, and find more polygons

## transformations in 2D

- coordinates
  - need point of origin (0,0) and axes (x and y)

- we want to define transformations generally, without need for coordinates
  - but hardware uses coordinates, so we must use them eventually
- dot product
  - product of magnitudes and cosine of angle between
    - \* or sum of product of coordinates along each axes
  - when dot product is 0, vectors are perpendicular
- 2d cross product
  - the cross product we normally think of only makes sense in 3d
  - our 2d cross product is just a vector of same magnitude, perpendicular to original
  - unary operation
  - represented by  $v$  superscript perpendicular-sign
  - $vp = (-v.y, v.x)$
  - $v$  dot product with ( $v$  cross product)  $== 0$
- there are two kinds of transformations
  - conformal:
    - \* preserves angles
    - \* translation, rotation, uniform scaling
  - affine (aff-in)
    - \* can be represented by multiplication by some matrix
    - \* translation, rotation, uniform/non-uniform scaling, shear
  - some conformal transforms are affine, but not all
    - \* conformal is affine if it can be represented as matrix
- translation
  - add a vector to every point
- uniform scaling
  - scale about a point (about an origin) by a scale factor
  - the point (origin) about which you scale will be unaffected by the scaling
  - the farther something is from the point (origin), the more its position will change
- non-uniform scaling
  - same as uniform scaling, but you now have a vector that you're scaling along
  - so take the vector from transform-origin to point, find parallel to transform vector, and scale that
  - scaling along a vector is not the same as scaling along the x and y components of that vector separately
- rotation
  - $q$  = vector from transform-origin to  $p$
  - new point is transform-origin + linear combination of  $q$  and  $q$ -cross determined by sin and cos of theta
- shear
  - not the same as non-uniform scaling
  - move point in direction of  $v$ , proportional to distance to  $o$  perpendicular

- ular to v
- reflection
  - TODO do we need to know this?
- matrix representation
  - compact
  - allows multiple transforms to be composed to single matrix (efficient)
  - if you have 3 points and those 3 points after some transformation, you can solve for the transformation
    - \* result is matrix
      - means that you can only solve for affine transforms
    - \* TODO do we need to know how to solve that on the exam?
- TODO how much of the transformation equations do we need to know?

## fractals and iterated function systems

- affine transform fractal is defined by set of contractive transformations
- contractive transform: transform  $F$  is contractive if for any two compact sets  $X1, X2$ , the distance between them is less after transforming them
  - that is,  $D(F(X1), F(X2)) < D(X1, X2)$
- hausdroff distance:
  - if two sets are equal, their distance is 0
  - distance of a,b is same as distance b,a
  - hausdroff distance is the maximum distance of a point in one set to the closest point in the other
- attractor: shape that fractal approaches after a large (ideally infinite) number of iterations
  - if transforms are contractive, attractor is independent of starting point(s)
- fractal tennis:
  - algorithm to draw fractal by randomly applying transforms to the same point
    - \* but need to iterate point for a few hundred iterations first to get it into the attractor
  - resulting fractal is not perfect
  - can be made better by weighting fractal transform random choice by area
    - \* difficult to calculate the area of a transform
      - singular value decomposition
      - get eigenvecgtors (or eigenvalues?)
      - are of eigenvectors is relative area of transform
- condensation set: basically a thing you add in at every iteration
  - allows shape to build on itself
- fractal dimension
  - like spatial dimension, but for fractals



- $dim = -\log(\#transformations)/\log(scalefactor)$
- fractal curves can have infinite length but enclose finite surface area
  - and that's fine
  - fractal paint bucket would not work because paint atoms have finite size

## transformations in 3D

- very similar to transformations in 2d
- things that are the same:
  - dot product, translation, uniform/non-uniform scaling
- cross product
  - now is binary operator
  - produces third vector that is perpendicular to both input vectors
  - magnitude: product of magnitudes and sine of angle between
    - \* mag represents area of 4-sided polygon formed by the two vectors
  - uses special matrix called  $\_$  (underscore)
    - \*  $v \text{ cross } \_ =$  put components of  $v$  in special places
    - \* then  $(v \text{ cross } \_) \text{ cross } w == v \text{ cross } w$
- rotation
  - input: axis to rotate about (specified as point and unit vector), theta to rotate
  - thus you're rotating in the plane that is perpendicular to the axis
  - component of  $q$  parallel to axis does not change, perpendicular component is rotated (in plane)
- mirror image
  - the same as non-uniform scaling with  $a = -1$
  - reflect about plane formed by normal vector  $v$  and point  $o$
- orthogonal projection
  - flatten things straight down onto plane
- perspective transformation
  - flatten things onto plane, but as if seen by an observation point  $e$  (an eye)
  - not defined for vectors (depends on where the vector is how much it gets scaled by)
  - **not an affine transformation!**
  - therefore you need to use the bottom row of the matrix also
  - the final 3d location is found by taking the normal 3d output and dividing by the 4th element (a scalar)
- hierarchical animation
  - split body into components joined by joints
  - each joint has a transformation associated with it, so you can apply the transformations corresponding to what position of components you want, and then render that

- skeletal animation
  - skeleton inside model mesh has hierarchical animation stuff
  - every vertex on mesh has a list of weights of how it's position depends on transformations of bones
  - thus allowing mesh to deform like the skeleton does
- OpenGL matrices
  - view, model, projection, viewport
  - view: position the camera
  - model: position model in world
  - projection: flatten world into 2d plane
  - viewport: transform projection into window pixel coordinates
  - opengl uses ModelView matrix
    - \*  $ModelView = V^{-1}M = T^{-1}R^{-1}M$
    - \* viewer always views from origin
      - TODO looking down negative z axis?
    - \* then you push a matrix for the models so that they're positioned correctly

## color

- human eye
  - cornea, iris, lens, retina
  - center of focus is fovea
  - stuff not in fovea (center of focus) is out of focus
    - \* this is different than blurry
    - \* all of the computer screen is in focus, so it can't perfectly replicate out-of-focus things
  - rods
    - \* brightness only
    - \* best response to blue-green light
    - \* mostly in peripheral, not fovea
    - \* more common than cones (100 million in retina)
  - cones
    - \* captures color
    - \* mostly in fovea, few in peripheral
    - \* far fewer overall than cones (6 million)
  - also at the very middle of your eye, where the optic nerve connects, there is no rods or cones
    - \* called blind spot
    - \* brain smooths this out
  - human vision
    - \* center of focus is highly detailed and in color
    - \* peripheral is black and white and less detail
    - \* in general, we can distinguish change in intensity more than

change in color

- intensity/luminance: how much light there is
  - like energy
- brightness: perceived intensity
  - color-dependent: because the rods respond differently to different wavelengths
    - \* not as much due to intensity dependence of cones
  - human eye can notice about 1% change in intensity
  - eyes more easily notice ratio between intensities than absolute intensities
    - \* so changes 1->2, 2->4, 4->8 all look about the same
    - \* to get equal-looking brightness increments, you need a power series:
      - minimum:  $I_0$
      - maximum: 1.0
      - series:  $I_0, rI_0, r^2I_0 \dots r^n I_0 = 1.0$
      - $n = -\log_{1.01}(I_0)$
- gamma correction
  - correct for how humans perceive color
  - combining colors is not linear!
    - \* you need to convert to linear space, mix colors, then convert back
  - gamma correction is meant to model old CRT displays
    - \* new displays do the same thing just for compatibility
- dynamic range
  - dynamic range =  $1/(\text{max intensity})$ , where minimum intensity == 1
  - different than contrast
- contrast
  - maximum vs minimum brightness a display can do *at the same time*

## coloring with limited intensities

- most displays are limited to 256 intensities (8 bits per channel)
- thresholding
  - naive approach
  - just round to nearest integer
  - does not usually look good, but is fast
- halftone
  - eyes integrate over area, so get varying intensity by having different intensities in an area
  - split image into blocks of pixels
    - \* e.g. like 4-pixel blocks
    - \* will lose resolution!
    - \*  $n * n$  block =>  $n^2 + 1$  intensity levels (+1 because includes both endpoints of all-on and all-off)
  - you then make one pattern per each intensity level

- \* assign based on average intensity level of block
- patterns
  - \* brain will recognize any pattern that exists
  - \* so make blocks random and uncorrelated so there is no pattern
- dithering
  - like halftone but we don't want to lose resolution
  - instead of using pattern to fill blocks, use pattern as threshold
  - fill in a pixel iff its intensity is greater than the threshold
  - preserves more fine details than halftone does (thanks to not losing resolution)
- error diffusion
  - visit pixels in a specific order (e.g. scanline order)
  - each time you round a pixel value, propagate that pixel's round-off-error to the adjacent pixels
    - \* but only pixels that have not been visited yet
  - can be combined with halftone/dithering
    - \* diffuse errors on block-by-block basis to next block
  - looks better than halftone/dithering

## color/light

- if all light entering eye is one wavelength, we see that color
  - but light is usually a spectrum of many colors
- receptor response
  - eye has 3 kinds of cones that respond to different kinds of light differently
  - not exactly one wavelength for each kind of cone
    - \* not actually all that close to RGB either
    - \* most response to yellow/green
  - just need to stimulate these 3 kinds of cones in the same way some wavelength of light would
    - \* don't actually need to produce that wavelength of light
  - the 3 types of cones in the eye makes the color space 3d
- CIE XYZ system
  - X,Y,Z are 3 primary colors. All colors can be made of linear combination of these
    - \* 3d color space
    - \* determine xyz coordinates for color using color matching function (using integration)
  - for all visible colors, x,y,z are positive
  - x,y,z are **not** visible colors by themselves!
  - visible light forms cone-ish shape pointing out from origin
  - luminance
    - \* intensity
    - \* 1-dimensional scalar

- \*  $x + y + z$
  - chromaticity
  - \* 2d quantity
- chromaticity diagram
  - is the  $x + y + z = 1$  plane for visible light
  - spectral colors
    - \* colors of rainbow, correspond to real wavelength of light
    - \* along the top curve of diagram
  - non-spectral colors
    - \* colors along the bottom edge of diagram
    - \* do not correspond to real single wavelength of light (but we still perceive them)
  - saturation: distance from white center point of diagram
  - hue: direction that from white center point
    - \* hue + saturation are another way to describe color
    - \* AKA dominant wavelength
  - complimentary colors
    - \* two colors that sum to white
    - \* e.g. white is halfway between them on the diagram
  - combining 2 colors
    - \* if you can produce two colors on the diagram, you can vary the intensity to get any two colors on the straight line between them
  - combining 3 colors
    - \* same as 2 colors, but you can now make any color inside the triangle they form
    - \* triangle forms a gamut
- gamut
  - range of colors that a device can make
  - represented as triangle on the chromaticity diagram
  - red, green, blue (RGB) allows you to cover most of the visible spectrum
    - \* nowhere near all though
  - different devices have different devices, so the same image might look different
    - \* need calibration

## color models

- RGB
  - red, green, blue
  - additive system
  - typical for monitors, because tells you what value for each pixel to use
- CMY
  - cyan, magenta, yellow
  - used in printing

- subtractive
  - complimentary to RGB:  $CMY = (1, 1, 1) - RGB$
- CMYK
  - cyan, magenta, yellow, black
  - best for printing, because you mostly print black
  - $K = \min(c, m, y)$
  - $C = C - K$
  - $M = M - K$
  - $Y = Y - K$
  - use as much black as you can, because black is cheap
- YIQ/YUV
  - NTSC, PAL
  - backward compatible with black and white because intensity is completely separate
  - also luminance is given more bandwidth because more important to eye
- HSV
  - hue, saturation, value
  - user-friendly way to specify color
  - value: like lightness of color
- Lab and Luv
  - perceptual-based model for color
  - not perfect because perception is not uniform among humans
  - better than RGB or XYZ though
- color representation
  - usually some number of bits per color channel
  - alternative: color indexing
    - \* store all colors in a table and index into the table for each pixel
    - \* good for limited palette
    - \* can then use dithering stuff to get more detail

## lighting

- global illumination
  - light from all sources, no matter how many times it had to bounce around the environment to get to your eye
- local illumination
  - light must go directly from light source to object to eye
- when calculating lighting, basically everything is a unit vector
- reflection models
  - how light interacts with a surface (without changing frequency)
  - ideal specular
    - \* reflection law
    - \* metallic, mirror

- ideal diffuse
    - \* lambert's law
    - \* matte
  - specular
    - \* directional diffuse
    - \* glossy
- illumination model
  - ambient, diffuse, specular
  - each of these is really 3 equations, one for each color channel (RGB)
- ambient
  - $I = k_a * A$ 
    - \*  $k_a$  = ambient reflection coefficient (material specific)
    - \*  $A$  = intensity of ambient light (constant)
  - uniform light caused by secondary reflections
  - lights up everything the same
  - accounts for all indirect illumination
    - \* which means it determines the color of the shadows
- diffuse
  - light reflects equally in all directions
  - $I = C * k_d * \cos(\theta)$ 
    - \*  $I$  = intensity
    - \*  $k_d$  = diffuse reflection coefficient
    - \*  $\cos(\theta) = \mathbf{L} \cdot \mathbf{N}$ 
      - $\theta$  = angle between (vector pointing from point to light source) and surface normal
    - \*  $\mathbf{N}$  = normal to surface
  - if light source is infinitely far away, you can assume the vector pointing toward it is constant. Otherwise, you must calculate the vector pointing toward it for every point on the surface.
- lambert's law
  - brightness of surface depends only on the angle that the light source makes with the surface
    - \* specifically, depends on cosine of angle that light vector makes with surface normal
    - \* or when they're both unit vectors, just use dot product
  - does not depend on viewer angle
  - thus, the object appears the same no matter where you view it from, as long as the object and light source do not change location
  - also implies light is reflected equally from the point of intersection in every direction
    - \* makes for a matte appearance
- specular
  - mirror-like reflection. forms highlights on shiny objects
  - $I = C * k_s * \cos^n(\alpha) = C * k_s * (\mathbf{R} \cdot \mathbf{E})^n$ 
    - \*  $\mathbf{R}$  = direction of reflection
    - \*  $\mathbf{E}$  = vector pointing toward eye

- \*  $n$  = specular component (controls size of highlights)
  - calculate  $R$ :
    - \* flip  $L$  around  $N$
    - \*  $R = 2(L \cdot N) * N - L$
  - finding reflected vector
  - $n$  exponent
    - \* larger  $n \Rightarrow$  smaller reflection thing
- multiple sources
  - only ever one ambient term
  - diffuse and specular are per-source
    - \* light is additive, so add light
- attenuation
  - light decreases the further from the source you are
  - various formulas to compute how much it attenuates
    - \* impacts performance depending on how complicated function is
- spot lights
  - light shines in a specific direction
  - so when you calculate light, you need to take the angle into account
- implementation considerations
  - if angle greater than 90, it's on the backside of the object
  - you need to negate the normal vector then
  - if you're only considering one-sided light, ignore the backside of the object
  - you want  $k_a + k_d + k_s \leq 1$ , or else you get saturation of colors
- OpenGL
  - set normals for surface
    - \* using `glNormal()`. Set until changed (so can be for multiple vertexes)
  - create/position lights
    - \* you can have different ambient/diffuse/specular for each light
    - \* must enable each light manually, and enable lighting in general
  - specify material properties
    - \* ambient, diffuse, specular
    - \* shiny is the  $n$  exponent
  - select lighting model
    - \* is viewer local?
    - \* is light two sided?