

# ECEN454 refsheets

<https://github.com/joshua-wright>

## specifically for final

- exam 1 was on March 8 (2017.03.08)
- nothing specifically from previous exam

## Propagation vs Contamination delay

- contamination delay ( $t_{cd}$ ): time from input change to **any** output changing value
- kind of a best-case delay or smallest possible delay
- time counted when input crosses 50% of logical high voltage level
- propagation delay: time from when all inputs are stable to when all outputs are stable
- kind of like a worst-case delay

## static vs dynamic power

- static power
  - consumed when circuit is not active
  - e.g. leakage current
  - static because no inputs are changing
- dynamic power
  - power consumed due to gates switching and stuff
  - e.g. charging/discharging parasitic caps, etc...

## Latch vs Flip Flop

- Latch is level-sensitive, flip-flop is edge-sensitive
- register is flip-flop # Latch/flip flop designs
- pass transistor latch
  - pros: tiny, low clock load
  - con: voltage drop across is  $V_t$ , it's non-restoring
  - con: not really a latch because doesn't hold signal
- transmission gate
  - pro: no  $V_t$  drop
  - con: requires inverted clock
- inverting buffer
  - TODO
- tristate feedback
  - first complete latch here
  - risk of backdriving: downstream could change state if it is very strong
- buffered output
  - no backdriving problem

- widely used
- rather large and slow though, and large load on clock
- a flip-flop is just two latches back-to-back

## Metastability

- stable works if it's not at a strong high or low but will settle somewhere
- metastable is where it might sit there, but if it must settle, you don't know which one it will settle to when perturbed

2017.03.27

## Sequential Circuits

- sequential means that the circuit holds state: the output depends on both current and past input
- to model a state machine, you can unroll it to [register]  $\rightarrow$  [combinational]  $\rightarrow \dots$ 
  - this way you can use regular timing stuff for it: arrival time, slack, etc...
- Mealy FSM: output of circuit is from combinational logic
- Moore FSM: output is from registers
- pipelined circuit: uses registers to hold state between clock cycles, because not all of the combinational logic can happen fast enough to work in a single clock cycle
  - pipelined circuits can use flip-flops or latches?
- the clock consumes 20-30% of the power on a chip
- the whole reason that registers are needed is because data (signals) moves through components at non-constant speed
- reset
  - can be sync or async
  - force low output when reset is high
- sequencing: it is (generally) equivalent to split the sequential logic in two, and then use two latches (one halfway through) instead of one large section of sequential logic with flip-flops at either end
  - this is called two phase clocking
  - you have to make sure the middle latch operates on clock-bar instead of clock

symbol	definition
$t_{pd}$	logic propagation delay
$t_{cd}$	logic contamination delay
$t_{pcq}$	Clk $\rightarrow$ Q propagation delay
$t_{ccq}$	Clk $\rightarrow$ Q contamination delay
$t_{pdq}$	D $\rightarrow$ Q propagation delay
$t_{setup}$	setup time of flip-flop/latch
$t_{hold}$	hold time of flip-flop/latch

- timing:
  - D $\rightarrow$ Q delay only makes sense for latches, since for flip-flops it is simply one clock cycle
- sequencing overhead and max delay:
  - $T_c$ : cycle time
  - need for combinational logic to be fast enough
  - for single phase flip flop:  $t_{pd} < T_c - (T_{setup} + T_{pcq})$
  - for two-phase latch:  $t_{pd} = t_{pd1} + t_{pd2} < T_c - 2t_{pdq}$
  - max delay is dictated by cycle time and sequencing overhead; sequencing overhead does not effect minimum delay
- minimum delay:
  - for single phase flip flop:  $t_{cd} > t_{hold} - t_{ccq}$

- \* combinational stuff must be slow enough that it doesn't violate hold time of second flip flop
- for two-phase latch:  $t_{cd1}, t_{cd2} > t_{hold} - t_{ccq} - t_{non-overlap}$ 
  - \* hold time applies twice each cycle, once per each combinational circuit
  - \*  $t_{non-overlap}$ : phase between clock signals?
- time borrowing: for two-phase latch-based system, you can borrow time from one segment to the other while the latch is transparent
  - works (is possible) as long as the total circuit still completes within one clock cycle
- clock skew:
  - when uncertain: causes tighter required CL propagation delay and minimum CL contamination delay, and reduces opportunity for time borrowing.
  - when know constant skew: can help with same effect as time borrowing, except can also have effect for flip-flop based systems
- skew tolerance:
  - flip-flop circuits are not very skew tolerant because the setup and hold times define a specific window of time in which the data must arrive
  - latch circuits are more hold-time-tolerant because the data can arrive any time that the latch is transparent
- if setup times are violated, you can reduce clock speed to fix it
- if hold times are violated, it won't work at any clock speed

2017.03.29

## Clocking

- H-Tree (clock distribution network)
  - laid out in such a way that the delay from the center to any leaf node is equal
- Grid clock distribution
  - grid on two or more levels
  - ensures low local skew, but there could still be larger chip-wide skew
- generally, the longer the delay between two point, the larger the worst-case clock skew
- Domino circuit: TODO
  - a kind of dynamic logic
  - when the clock pulse cascades through the circuit, activating the pMOS at each stage?

2017.04.03 (+more)

## Memory

- random access or sequential access, or content addressable
- CAM: content-addressable memory
  - basically like an associative array, or database
  - designed to search all memory at once in a single operation (wikipedia)
  - used in caches
- SRAM:
  - 12T and 6T cell designs
  - 12T design: just a latch connected to a bitline
    - \* boring and impractical because it's too large
  - 6T design: cross-coupled inverters (3T each) +2T for bit and bit-bar lines
    - \* we focus on 6T design
  - read: pre-charge bit and bit-bar, raise wordline, read result
    - \* (need to pre-charge or else read is destructive?)
    - \* need sense amplifiers at the end of the bitlines because the signal will be weak, and the parasitic capacitance of the bitline depends on how many cells are 0 or 1
  - write: drive data onto bit and bit-bar and raise wordline to write
  - cross coupled inverters must be much larger than the bitline transistors so that it doesn't flip while reading

- per each column, you need:
  - \* piece of decoder
  - \* bitline sense amplifier
  - \* column multiplexer
- multi-port: means you can read from and write to the same cell in the same clock cycle and it works
  - \* implement by having multiple sequential read/write per cycle?
- decoders:
  - must be pitch-matched to SRAM cell width for layout efficiency/simplicity (requires very skinny gates)
  - large ( $n > 4$ ) decoders are inefficient, thus it is often helpful to use predecoding: factor out common gates into a second stage of decoders.
    - \* smaller area as alternative, but same logical effort (since same number of gates traversed)
- twisted bitlines: bitlines are long and right next to each other, so there will be significant crosstalk and parasitic cap
  - solution: swap adjacent bit and bit-bar lines occasionally
  - reduces Miller factor (MCF)
- DRAM:
  - one transistor and one capacitor
    - \* transistor gates capacitor
  - no charge is 0, charge is 1
  - open transistor to read/write
  - read is destructive
  - must be refreshed periodically due to leakage
    - \* thus, it's always consuming power
- shift register:
  - can be implemented with cascade of flip-flops
    - \* can lead to hold time violations because no delay from clock to next data line? (TODO what? how?)
    - \* not very dense
  - store data in SRAM and store begin/end pointers
    - \* more dense than individual flip-flops
- Tapped Delay Line
  - shift register with programmable number of stages
  - implement using sequence of flip-flops, each with a bypass MUX from data in to out
- Queues: are a thing (TODO)
- ROM: Read Only Memory:
  - mask-programmed ROMs are one transistor per bit
- PROMs, EPROMs, etc
  - generally burn out specific fuses to program ROM
  - E for erasable

2017.04.10, 2017.04.12

## Low Power Design

- not so important when running from a battery since batteries store energy, not power
- peak power
  - determines power/ground wiring and packaging limits
  - impacts signal/noise ratio
- reduce dynamic power: (AKA active power)
  - use smaller transistors
  - clock gating
  - reduce  $V_{DD}$
  - reduce frequency
    - \* but be careful because that will make the computations take longer
- reduce static power
  - power gating

- selectively use lower  $V_t$  devices
- stacked devices
- body bias
- clock gating: disable the clock signal to an expensive component
  - eliminates that component's dynamic power (but of course you can't use it)
  - doesn't change static power
- voltage islands: different  $V_t$  for different devices
  - allow low  $V_t$  devices to run slower because they don't need to be fast (maybe they're not in the critical path)
  - often different voltage levels are alternated row by row in standard-cell layouts
  - you must convert from one logic level to another, because high/low will be different
    - \* sometimes it's possible to avoid conversion, e.g. by using a buffer that you would have anyway, with inverters built from half high- $V_t$  transistors and half low- $V_t$  transistors
  - there are all kinds of tricks for determining which components should be high/low voltage and where to put voltage islands and level converters
- DFVS: Dynamic Frequency and Voltage Scaling
  - requires:
    - \* programmable clock generator (PLL)
    - \* supply regulation loop to for setting minimum  $V_{DD}$
    - \* software support for deciding best stable frequency and voltage for given load and performance goals
  - voltage regulation loop
    - \* voltage divider would be horribly inefficient
    - \* instead, rapidly charge/discharge capacitor
- stacked devices:
  - stacking devices reduces leakage
  - because leakage is exponentially proportional to  $V_{ds}$
- body bias AKA dual  $V_t$ 
  - changing  $V_t$  at runtime is called Adaptive Body-Biasing (ABB) or Dynamic Threshold Scaling (DTS)
    - \* requires triple well fabrication process, because what would have been lowest level substrate must itself sit inside a well that will be biased
  - design with low and high  $V_t$
  - low  $V_t$  is faster than high  $V_t$  but consumes 10x more power
    - \* because reducing  $V_t$  increases sub-threshold leakage exponentially
    - \* but reducing  $V_t$  decreases gate delay (increasing performance)
  - to increase  $V_t$ , put a negative bias on  $V_{SB}$  of nMOS
    - \* use DC charge pump to get desired voltage
- power gating: use sleep transistors
  - basically cut off  $V_{DD}$  from power-hungry circuits when not in use
  - in sleep mode, the sleeping transistors are stacked with the transistors that are causing them to sleep, thus reducing leakage even more
- MTCMOS sleep transistors
  - combines power gating and dual  $V_t$
  - TODO what is this?

2017.04.26

## Testing